

How does an Agent use Slack?



Slack workflows



Automatic result checks

- I studied how agents observe and act on pages, then reviewed Slack mockups and agent environments.
- I mapped search, threads, and message edits to agent actions and saved changes I could verify.
- I built the app, action loop, and result checks separately so I could test each part.



Research
computer use



Map Slack
workflows




Build the
environment



Test and
record runs

I started with a task that needs a **real saved change**

- The agent must find Alex's launch-review message and change its time from 14:00 to 15:00 UTC.
- Typing the new time is not enough: it must save the edit without changing other messages.



- A static mockup cannot test persistence. A full Slack backend adds services these tasks do not need.
- Fictional data · Task abbreviated here · [Try the workspace without a key](#) · [Prior-work review](#).

The task set covers 18 workflows

Six focused tasks

Update a topic

Channel → edit topic → save

Reply in a thread

Find request → open thread → reply

Edit a message

Find own message → edit → save

Triage an incident

Find active incident → react → pin

Delete a draft

Find own draft → delete → confirm

Send a handoff

Read current details → compose DM

Coordinate work

Coordinate a release *

Read sources → edit → reply → topic → DM

Close out an incident

Read resolution → replace pin → reply → topic

Refresh saved work

Find current sources → replace saved items → DM

Record a design decision

Read approval → react → pin → save → details

Repair an outdated handoff

Read correction → edit DM → react → save

Complete a QA sign-off

Read test results → reply → fix reaction → save

Combine several changes

Publish a release update

Read source → format update → pin → delete draft

Brief the on-call engineer

Read rotation → DM primary → save → details

Correct a threaded estimate

Read review → edit own reply → react → reply

Refresh the launch board

Find plan → replace pin → save → topic

Hand off an approved design *


Read decision → react → DM → reply → save

Prepare a release retrospective

Read three sources → summarize → pin → save → DM

- These diagrams summarize the intended workflows. Exact instructions and preservation checks remain in each task.
- * Two original tasks contain wording problems. Their results remain unchanged and are discussed later.

The model proposes an action; the harness executes it



- The harness accepts only supported actions, so a model response cannot execute arbitrary code.
 - After Save, it reads the updated page so the model can check whether the edit appeared.
 - Action, time, and spending limits stop a stuck agent from making model calls indefinitely.
-
- The example shows browser control. The API path executes application operations instead.

The model can use **four interfaces** to the same app

Interface	Observation → action	Interaction
 Pixels	Screenshot → mouse and keyboard	Browser
 Accessibility	Named page controls → click or fill	Browser
 Page JSON	Structured page controls → click or fill	Browser
 Actor API	Workspace data → direct changes	Tool use

- A browser edit requires finding the editor, entering text, and saving the change.
- The actor API can request that edit directly, so the workspace changes without a browser click.
- The same grader checks the result, but the model uses different information and actions to reach it.

-
- Page JSON describes page controls, not the underlying database. Actor API access exposes broader structured data.
 - The harness uses a shared JSON action protocol, not each provider's native computer-use protocol.

Success means the requested change was saved

- A model can type the right text and click Cancel, so the screen alone cannot prove success.
- The grader is code that checks the stored messages, topics, and other workspace data.



Required saved message

“Launch review is at 15:00 UTC. Please bring the final checklist.”



Right change

Exact text in Alex's original message



No extra changes

No new message or unrelated edit

Both pass

Reward = 1

- The grader checks the final state. It does not penalize a wrong action that the agent fully undoes.
- Task definitions and result checks.

I tested the success rule on the same edit

- Alex's original message must change from 14:00 to 15:00, with no other changes.

APP TEST

Save the requested edit

A script uses the UI to edit Alex's message, then reads the stored data.

15:00 saved → expected pass

GRADER TEST

Give it incorrect saved data

Keep 14:00, or save 15:00 and add an unrelated message.

Either case → expected failure

RECORDING TEST

Recheck the saved run

Run the same grader on the archived final workspace, without calling the model.

Recomputed grade = recorded grade

- This helps distinguish software errors from a model taking the wrong action.
- Passing these tests does not prove the task instructions are clear.

-
- [Software verification](#) · Broader workflow tests cover 84 valid and 2,583 invalid states across seven seeds.

Each run starts with its own copy of the task

- Even with one model, each attempt needs fresh data so it cannot inherit completed work.
- Each run gets its own SQLite database and browser session, so its actions change only its own workspace.

Same task and starting data • separate copies

RUN A

Browser session A → database A

Agent saves its edit

Launch review at 15:00

RUN B

Browser session B → database B


No action in this run

Launch review stays at 14:00

The new run starts with the message at 14:00. The previous result stays in history.

- Tested: write to A, check B is unchanged; close A, check B still works. Cost: separate files and browser setup.
- The actor and control servers share a process. Separate data is not isolation from host crashes or malicious code.

What makes one run take longer?



- A fresh workspace makes runs independent, but each extra decision adds another model request.
- In the live viewer, the workspace waits during the model request and updates after the action executes.

-
- A faster reset helps between runs, not between decisions. Local M5 Max / Node 24 medians exclude model calls and are not a full-run estimate.
 - Samples and setup: 100 stored sessions; screenshots from 10 sequential browser contexts, not 100 concurrent agents.

Each run consumes **memory and storage**

- The browser stays open while the model decides, so a slow run also holds memory for longer.

WORKSPACE DATA

One SQLite file per run

Changes stay inside this attempt

48 KiB per fixture after one post

BROWSER EXECUTION

Browser stays open

Used for actions, observations and the live view

Uses memory until the run closes

EVIDENCE

Recordings remain

Review and replay without new model calls

Disk use continues after cleanup


- More simultaneous runs need more browser memory.
- More saved runs need more storage.

-
- 48 KiB is fixture disk size, not memory per run. Full browser memory and recording size need separate measurements.
 - [Resource choices and measurements](#) • Experiment runs close their own browsers. Browser reuse is untested.

A queue would limit how many browsers run at once

- Because each active run holds a browser in memory, starting every visitor immediately could exhaust the machine.
- I would queue new runs until a worker is free, while keeping each active run's data separate.

PROPOSED, NOT IMPLEMENTED OR LOAD-TESTED



A visitor would wait for a free worker before their run starts. More workers could reduce the queue, not model response time.

- Concurrent capacity is unmeasured. Test 1, 4, 8, and 16 workers for throughput, memory, provider limits, and crash cleanup.
- Use separate processes or containers for stronger containment. Browser pooling trades lower startup cost for shared failure risk.

The reusable harness made **model comparisons possible**

- Once runs were repeatable, I wanted to compare how different models completed the same tasks.
- I expanded the assignment into a benchmark prototype with two separate comparisons.
- The first compares models using accessibility controls, and the second compares four interfaces.

Original assignment



Working environment

Focused Slack workflows
Known starting state



Repeatable harness

Reset · act · check · record
Select a model



Expanded scope



Benchmark prototype

306 model-comparison runs
96 interface-comparison runs

The next slides cover the test conditions, results, and limitations of each comparison.

Comparison 1: 17 models, using accessibility controls

$$\begin{array}{ccccccc} \text{📁} & & \text{☑️} & & \text{🛡️} & & \\ 17 & \times & 18 & \times & 1 & = & 306 \\ \text{model routes} & & \text{tasks} & & \text{attempt per pair} & & \text{recorded attempts} \end{array}$$

- Each model receives named page controls and returns one JSON action per request.
- All 306 runs use the same starting data, recent-four history, and no site guide.
- This tests model performance under one interface, not the best interface for each model.

40 actions 180 s per run \$5 estimated allowance per run

- Seed 1042 • 4,096 output tokens per request • identical graders • provider-default reasoning.
- I kept 105 earlier runs and added 201 new runs without replacing failures.
- [Node 24 setup and API key instructions](#) • [Exact routes, limits, and run history](#). Provider routing and collection time were not controlled.

Results for 17 models

1 per task/model · Collection complete

306 / 306 attempted · 152 passed · 61 incomplete · 93 blocked

\$189.3110

Usage estimate

\$3.7794

Unresolved · 35 calls

\$193.0904

Total allowance · 306 trials

Model ↕	Trials ↕	Passed ↕	Incomplete ↕	Blocked ↕	Median time ↕	Usage estimate ↕	Unresolved ↕	Total allowance ↕
 GPT-6 Astra	18 / 18	14 / 18	1	3	27.1s	\$35.9623	\$0.0000	\$35.9623
 GPT-6.1 Sol	18 / 18	14 / 18	4	0	28.5s	\$8.4352	\$0.0000	\$8.4352
 GPT-6 Luna	18 / 18	9 / 18	8	1	39.4s	\$0.5394	\$0.0092	\$0.5486
 GPT-OSS 120B	18 / 18	5 / 18	2	11	180.1s	\$0.6682	\$0.1358	\$0.8040
 Claude Fable 5.1	18 / 18	13 / 18	1	4	62.8s	\$48.4714	\$0.0000	\$48.4714
 Claude Opus 5.5	18 / 18	13 / 18	4	1	38.3s	\$28.6987	\$0.3741	\$29.0728
 Claude Sonnet 5.5	18 / 18	7 / 18	4	7	39.9s	\$13.9323	\$0.3304	\$14.2627
 Claude Haiku 4.5	18 / 18	6 / 18	3	9	28.3s	\$3.7921	\$0.0000	\$3.7921
 Grok 4.7	18 / 18	10 / 18	6	2	34.3s	\$11.9101	\$0.3809	\$12.2910
 Qwen 3.8 Max	18 / 18	7 / 18	1	10	125.1s	\$6.9999	\$1.0639	\$8.0638
 DeepSeek V4 Pro	18 / 18	8 / 18	1	9	106.8s	\$4.6816	\$0.2547	\$4.9363
 DeepSeek V4.1 Flash	18 / 18	7 / 18	3	8	65.9s	\$1.3825	\$0.0541	\$1.4365
 GLM 5.3	18 / 18	7 / 18	0	11	125.3s	\$3.6307	\$0.3453	\$3.9760
 GLM 5.3 Flash	18 / 18	8 / 18	0	10	46.2s	\$0.4444	\$0.0000	\$0.4444
 Kimi K3	18 / 18	11 / 18	3	4	50.0s	\$14.0051	\$0.7958	\$14.8009
 MiniMax M3	18 / 18	5 / 18	10	3	25.9s	\$1.4698	\$0.0353	\$1.5052
 Nemotron 3 Ultra	18 / 18	8 / 18	10	0	50.1s	\$4.2874	\$0.0000	\$4.2874

Passed Incomplete Blocked Not run

- 18 public tasks. 1 run per task and model. All use accessibility controls. Median time includes blocked runs.
- Total allowance adds the usage estimate to money reserved for missing usage. These are estimates, not invoices. Including earlier tests: \$194.1058 of the \$300 budget. [Full costs CSV](#) · [JSON](#)
- 105 earlier runs are included. No retries. One run per task cannot establish a reliable model ranking. Two tasks have unclear instructions. [Read the failure analysis.](#)
- [Test plan](#) · [Recorded runs](#) · [Download all trials \(CSV\)](#)

Why was the highest score 14/18?

- GPT-6.1 Sol and GPT-6 Astra each passed 14 tasks and missed the same four.

Design handoff	Both sent DESIGN, as quoted in the task. The grader expected Willow, but the task never said to replace DESIGN.
Release synchronization	The source location was ambiguous. Sol found the facts but left required changes unfinished, so wording alone does not explain the miss.
Release retrospective	All 17 models hit a limit before finishing. These attempts do not show whether a larger limit would be enough to complete the task.
Handoff repair	Both entered the correct edit, then repeatedly clicked Cancel instead of Save changes. Five other models passed.

- On the last three tasks, Sol reached 40 actions and Astra hit the \$5 estimated-allowance limit. That is not billed cost.
- [Failure analysis and traces](#). Across all models, 15/18 tasks passed at least once. One attempt per model and task; no grades changed.

What the early interface tests taught me

- Early interface tests recorded successful, incorrect and interrupted runs.
- They used different task and model combinations, so their pooled scores could not compare interfaces.
- For the next study, I gave every model the same six tasks through all four interfaces.



Later study: one model + one task

Four fresh workspaces • identical starting data

 Accessibility

{ } Page JSON

 Pixels

 API • tool use



Same final-state checks • 4 models × 6 tasks × 4 interfaces = 96 attempts

- Early development traces are separate evidence, not a third benchmark or part of the later scores.
- Later study: randomized interface order • fixed seed, history, reasoning, and limits • \$25 total • no retries. Full design and controls.

Interface comparison: 96 recorded attempts

Passed / attempted · median elapsed time in seconds · six tasks per cell

Model ↕	Accessibility ↕	Page JSON ↕	Pixels ↕	Actor API ↕	Allowance ↕
 GPT-6.1 Sol	<u>6 / 6</u> 0 incomplete · 0 blocked All 28.4 s Passed 28.4 s	<u>5 / 6</u> 0 incomplete · 1 blocked All 33.2 s Passed 50.3 s	<u>1 / 6</u> 0 incomplete · 5 blocked All 107.4 s Passed 76.7 s	<u>6 / 6</u> 0 incomplete · 0 blocked All 10.0 s Passed 10.0 s	\$3.9430
 Claude Sonnet 5.5	<u>4 / 6</u> 0 incomplete · 2 blocked All 26.3 s Passed 10.2 s	<u>4 / 6</u> 1 incomplete · 1 blocked All 26.5 s Passed 10.8 s	<u>2 / 6</u> 1 incomplete · 3 blocked All 23.7 s Passed 19.2 s	<u>5 / 6</u> 1 incomplete · 0 blocked All 9.9 s Passed 7.8 s	\$7.2784
 Qwen 3.8 Max	<u>5 / 6</u> 1 incomplete · 0 blocked All 26.7 s Passed 14.7 s	<u>4 / 6</u> 1 incomplete · 1 blocked All 23.9 s Passed 36.6 s	<u>0 / 6</u> 0 incomplete · 6 blocked All 55.2 s Passed N/A	<u>5 / 6</u> 0 incomplete · 1 blocked All 11.9 s Passed 6.3 s	\$5.4829
 Grok 4.7	<u>3 / 6</u> 1 incomplete · 2 blocked All 20.8 s Passed 7.8 s	<u>3 / 6</u> 3 incomplete · 0 blocked All 17.2 s Passed 7.8 s	<u>0 / 6</u> 1 incomplete · 5 blocked All 40.5 s Passed N/A	<u>5 / 6</u> 1 incomplete · 0 blocked All 9.0 s Passed 5.2 s	\$6.0573

[All traces and replays ↗](#) [Costs CSV ↓](#) [Paired results and timing ↗](#)

- All-attempt medians include early stops, while passed-run medians use only successful tasks.
- Passed tasks differ across cells, so these medians do not establish a speed ranking.

- Verified: 58 passed · 11 incomplete · 27 blocked. Usage estimate \$20.7180 + unresolved \$2.0435 = \$22.7616 of \$25.
- N/A means no passed task. Time includes model calls and browser work, not just model latency. Click a pass count for traces and replays.

Do the interface results hold across repeated attempts?

The first study used one attempt per condition. This follow-up tested each condition twice.

Passed / 6 attempts · three tasks, two repeats · median elapsed seconds


Model ↕	Accessibility ↕	Page JSON ↕	Pixels ↕	API ↕	Allowance ↕
 GPT-6.1 Sol	6 / 6 0 incomplete · 0 blocked All 17.2 s Passed 17.2 s	5 / 6 0 incomplete · 1 blocked All 23.1 s Passed 17.6 s	1 / 6 0 incomplete · 5 blocked All 180.2 s Passed 159.4 s	6 / 6 0 incomplete · 0 blocked All 8.7 s Passed 8.7 s	\$3.3918
 Claude Sonnet 5.5	4 / 6 0 incomplete · 2 blocked All 13.5 s Passed 11.4 s	4 / 6 0 incomplete · 2 blocked All 13.6 s Passed 9.9 s	3 / 6 1 incomplete · 2 blocked All 32.4 s Passed 14.1 s	6 / 6 0 incomplete · 0 blocked All 6.3 s Passed 6.3 s	\$5.5759

[All 48 traces and replays ↗](#) [Costs CSV ↓](#) [All paired comparisons ↗](#)

- Each interface had 12 attempts: API passed 12, accessibility 10, Page JSON 9, and pixels 4.
 - On 10 shared passes, API used a median 6.3 fewer seconds and 3 fewer actions than accessibility.
 - API skips browser interaction, so its higher completion count does not show better computer use.
-
- Verified: 35 passed · 1 incomplete · 12 blocked. \$8.2343 usage + \$0.7334 unresolved = \$8.9677 of \$25. Six requests lack receipts.
 - Two repeats each of thread reply, message editing and incident closeout. The earlier 96 attempts remain a separate study.
 - All-attempt times include stops. Shared-pass timing selects on success. Low-detail pixels, macOS keys and two repeats limit what this comparison establishes.

The next version needs validated, unseen task variants

- Ask someone outside the project to solve each task from its instructions.
- Compare their interpretation with the expected result.
- Set run limits with both short and long tasks.
- Use new task variants and repeat runs to measure reliability.



The current results describe these attempts. Repeated runs are needed to estimate reliability.

- [Detailed engineering review](#) · [Remaining validation work](#). These are next steps, not completed results.

Demo and design questions



01 · TRY

Use the workspace

Search, open a thread, edit



02 · COMPARE

Choose a model and task

All 306 attempts and costs



03 · REVIEW

Inspect the run

Actions, changes, result checks

- Which workflows should the next version test?
- Should the score include wrong actions that the agent later undoes?
- How many agents must run at once, and what should each run cost?

-
- Relay can reset, execute actions, and return a reward. No reinforcement learning training has been run.
 - [Code and setup](#) · [Requirements and remaining work](#).